Table 4: $\text{TPR}_{@\text{th}}$ of synthetic image detectors evaluated on laundered data. For completeness, we also report results for the original real images. Notice that, when considering only real images, this metric corresponds to the rate of false alarms produced by each detector. In bold, we highlight results lower than 30% on original images (which indicates that the detectors are correctly classifying them, on average). Only the detector [6] misclassifies real images, failing on 40% of them. Regarding laundered data, the detectors [3, 7, 4] identify the vast majority of SD-laundered samples as synthetic, except for [4], which behaves differently on SD 3.5 images. All other detectors do not substantially modify their performance on laundered data with respect to the original samples. In general, Flux-based laundering alters detection results much less than SD-based laundering. These findings further support our hypothesis that peaks are not crucial for deep-learning detectors. Conversely, linear detectors that explicitly rely on peaks are strongly affected by the laundering process.

Detector	Original	Laundered		
		SDXL	SD3.5	Flux1
[3]	**0.040**	0.976	0.985	0.361
[5]	**0.104**	0.203	0.116	0.127
[7]	**0.106**	0.979	0.811	0.436
[4]	**0.131**	0.862	0.328	0.206
[8]	**0.101**	0.187	0.132	0.118
[6]	0.403	0.393	0.406	0.406
[9]	**0.122**	0.167	0.270	0.176
Linear-8	**0.106**	0.822	0.929	0.808
Linear-16	**0.128**	0.756	0.881	0.742